

AGENT FLEET · FIELD NOTES

The 5 checks before an agent gets write access

I run a fleet of agents on a Raspberry Pi.
These are the gates I built *after* things broke.

THE WHOLE CHECKLIST

Five gates, in order

- Where do its secrets live?
- What can it reach?
- What does it do when it fails?
- Can it talk to anyone?
- How do you roll it back?

CHECK 1 / 5

Where do its secrets live?

I once pasted an API key into a chat message. It propagated to **eight** places before I caught it. Logs, configs, backups, a shell history.

One file, `chmod 600`. Never in the repo, never in a prompt.

CHECK 2 / 5

What can it reach?

A read-only agent is a tool.

An agent with write access is a **participant**.

Decide which one you are deploying, on purpose, before it runs.

CHECK 3 / 5

What does it do when it fails?

Silent failure is the expensive one.

An agent of mine filtered messages by its own name, matched nothing, and went quiet for hours. No error. No alert.

Failures must be **loud**. If it can fail without telling you, it will.

CHECK 4 / 5

Can it talk to anyone?

The moment an agent can message an outside party, it stops being automation and becomes a **spokesperson**.

Your tone, your claims, your liability.

CHECK 5 / 5

How do you roll it back?

If you cannot answer this in one command, you are not ready to give it write access.

Rollback is not a nice-to-have. It is the thing that makes the other four checks affordable.

THE SHORT VERSION

**Identity is what
turns
a tool into an
agent.**

FIELD NOTES FROM RUNNING AGENTS

And everything that makes agents useful
is also what makes them risky.